



# Ising Model

with Parallel Implementations

Taj Gillin




# 01


## Background

Ferromagnetism and the Ising Model

# Ferromagnets




Atoms have magnetic moments



Aligns with magnetic field below critical temperature


# Ferromagnets



At low temperatures, atoms align with each other and spontaneously polarize!

# Ising Model

x



+


- Spins on a lattice, each with  $\pm 1$
- Interaction: spins favor alignment with neighbors
- Simple but captures essence of phase transitions and ferromagnetism

x




# Ising Model

×



+



×


But in 4d!



# Monte Carlo (Metropolis) Algorithm


1. Randomly pick a spin
2. Compute energy change  $\Delta E$
3. If  $\Delta E < 0$ , flip
  - a. Else, flip with probability  $\exp(-\Delta E/kT)$
4. Repeat

$$\Delta E_i = -\sum_j J s_i s_j - h s_i$$



# 0

×




×

# 2 Approach

Non-parallel and parallel implementations



+







# Sweeping Method vs Random Updates

Instead of randomly picking spins, sweep over the whole lattice and visit every spin



Less randomness, but still has proper statistical sampling and achieves equilibrium

# Red-Black Updates

×



×




With sweeping, helps improve reliability and prevent pattern formation (see later)

×




# MPI Domain Decomposition

×




# MPI Domain Decomposition

x




# MPI Domain Decomposition



×


# MPI Domain Decomposition

×



- Decompose  $L \times L \times L$  lattice into sub-lattices, each handled by an MPI rank
- Each rank stores a portion of the lattice plus halo (ghost) cells
- Halo exchanges ensure each process sees correct neighbor spins


# Halo Exchange



- Neighbor ranks exchange boundary data (planes)
- Synchronized before next update

# GPU Kernel


- Each MPI rank individually runs kernels
- Spin updates, energy and magnetism are offloaded to GPU
- (1) Halo is exchanged on CPU, (2) data is copied to GPU, (3) calculations performed, (4) then edges copied back
- Operations seen in rocprof





0

×




×

# 3 Results

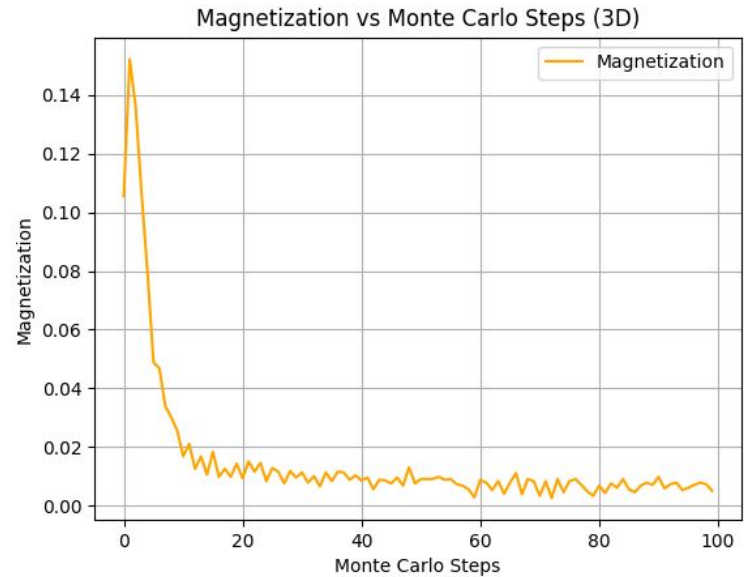
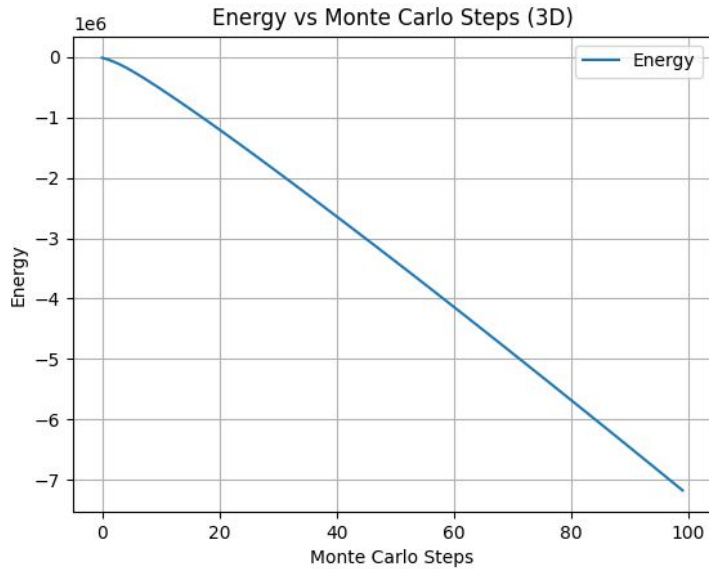
Issues and Speedup



+



# Serial Implementation



# Serial Implementation

x




<https://tajgilllin.neocities.org/cpu/even>

x



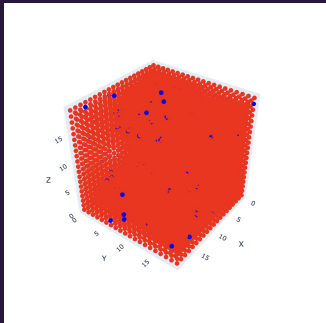

<https://tajgilllin.neocities.org/cpu/odd>

x



# Red Black

×




×




×


<https://tajgilllin.neocities.org/redblack/rb>





# Serial Implementation w Red Black



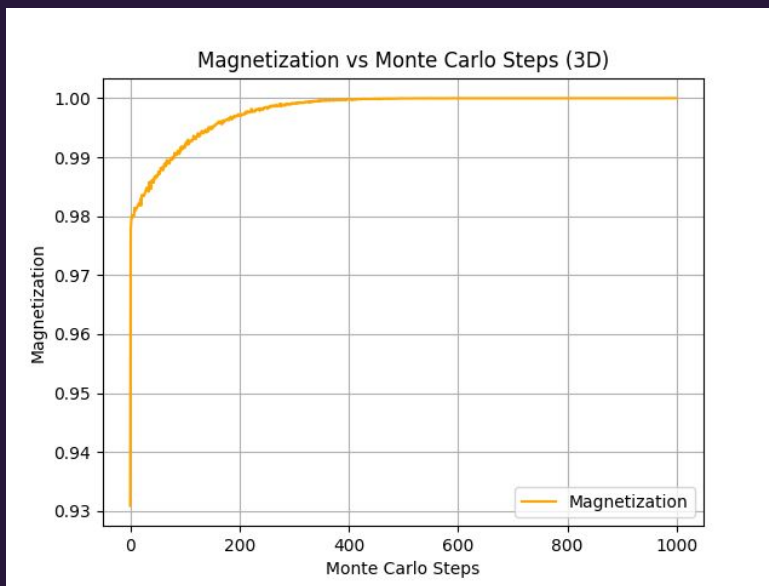
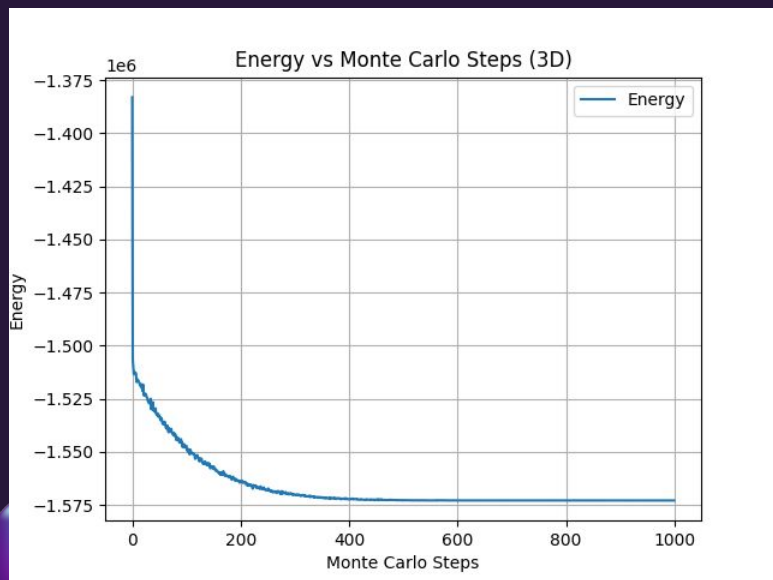
×



# MPI w Red Black



# MPI w GPU Kernels




# Comparison

64x64x64, 1000 steps,  $T = 2.5$ ,  $J = 1$

<b>Serial</b>	36.2 seconds
<b>MPI v1</b>	52 seconds
<b>MPI v2 (1 rank)</b>	35.6 seconds
<b>MPI v2 (8 ranks)</b>	5.8 seconds
<b>GPU</b>	0.95 seconds
<b>GPU (256<sup>3</sup>)</b>	28.6 seconds




# Roofline Model



CPU: MI250X HBM2e  
GPU: AMD EPYC 7V13


# Scalability

×




$64^3 \sim 262k$

×



$256^3 \sim 17m$




×

# Optimizations

- ◆ Shared GPU memory (0.95 → 0.71 seconds) for threads within block with shared memory
- ◆ GPU aware halo process, share directly from GPU to GPU (→ 0.28 seconds)
- ◆ With  $256^3$ , brought from 28.6 seconds to 5.6 seconds

×

×




×

# Future Directions

- ◆ Pinned memory (currently using device), faster transfers and communication
- ◆ Streams, compute edges, then start halo exchange while computing inside

×

×





# Thank you!